

The Algebraic Mind Revisited:

From Classic Critique to the Modern AI Debate

Jonas Peters, Philipp Haack

August 13, 2025

Contents

1	Introduction	3
2	Relations between Variables	5
2.1	UQOTOM	5
2.1.1	Humans can freely generalize UQOTOMs	7
2.1.2	Multilayer Perceptrons and Operations over Variables	7
2.2	Marcus's proposed Alternatives	13
2.2.1	Encoding Schemes	14
2.2.2	Case Studies	16
2.3	Summary	17
3	Contemporary Research	18
3.1	The Illusion of Thinking	18
3.2	A Case for Neuro-Symbolic AI: AlphaGeometry	21
4	Summary	23

1 Introduction

The human brain is a remarkably complex and efficient system, capable of performing an extraordinary range of cognitive tasks with ease. Its ability to learn from experience, adapt to new environments, and process vast amounts of sensory information has long fascinated scientists and engineers alike.

The study of its functionality is partially in the domain of computer science: We try to model the brain with computer algorithms, which is inspired by the idea that the human brain consists of elementary computational units, neurons, which itself can be sufficiently modeled with computers, and whose interplay creates the emergent complex behavior we see. This leads to the following premise:

Premise 1.0.1. *In theory, the human brain can be sufficiently modelled by a computer program.*

Over many decades scientists have been trying to make progress towards this monumental task of modelling the human brain. As a result, two paradigms on how to implement intelligent models emerged, they are called *connectionism* and *symbolism*. While connectionist models rely on training data, symbolist models have all their inference rules and algorithms directly coded into them. There also exist hybrid models, summarized under the term *neuro-symbolic AI*. Marcus claims:

Marcus (2024)

There can be no path to AGI without neurosymbolic AI.

To understand why Marcus is so critical of connectionist models which include LLMs, we will analyze one of his early writings, namely *The Algebraic Mind* Marcus, 2001. In this book, he argues that the popular architecture of MLPs (multi-layer perceptrons) is not sufficient to model the human brain, because it cannot represent certain algebraic structures. Even as of today, Marcus likes to refer to his early work, which emphasizes its importance:

Marcus (2024)

You would think the need for this reconciliation is obvious, and indeed it was the central point of my 2001 book *The Algebraic Mind*, which in the words of the subtitle sought to integrate connectionism (neural networks) with the cognitive science of symbol-manipulation:

In this essay, we will analyze Marcus's critique of connectionist models and his arguments against the MLP architecture, where we focus on chapter 3 *Relations between Variables*. We will also discuss the relevance of his critique in contemporary artificial intelligence research, especially in the context of LLMs.

2 Relations between Variables

Marcus starts with the following observation:

Observation 2.0.1. The human brain is capable of performing operations like addition, multiplication, inflecting verbs, and so on.

What do all these algebraic operations have in common? They all operate on variables based on rules. Thus, Marcus calls them *relations between variables*.

2.1 UQOTOM

Instead of considering all these functions of human cognition at once, Marcus focusses on a very special subclass of operations. He refers to them as *UQOTOM*, which stands for *universally quantified one-to-one mapping*. Like the name suggests, they are defined as follows:

Definition 2.1.1 (UQOTOM). A mapping $f : X \rightarrow Y$ is a UQOTOM if it is *one-to-one* and *universally quantified*. Universally quantified specifies that f is defined for all $x \in X$, i.e. that f is a valid function in the mathematical sense.

Some important examples of UQOTOMs that Marcus is going to analyze are the identity function, and the simple past operation which produces the simple past of a verb. According to Marcus:

Marcus (2001)

I do not mean to suggest that UQOTOM are the only mappings people compute. But UQOTOM are especially important to the arguments that follow because they are functions in which every new input has a new output. Because free generalization of UQOTOM would preclude memorization, evidence that people (or other organisms) can freely generalize UQOTOM would be particularly strong evidence in support of the thesis that people (or other organisms) can perform operations over variables.

Based on this quote, we can see that Marcus is interested in UQOTOMs because they enforce a system to *generalize* mappings. That is, given a never-before-seen input, a potential model implementing this UQOTOM must produce a new output that has not been produced before. And this generalization is a central aspect of *symbol manipulation* according to Marcus.

Marcus (2001)

To a system that can make use of algebra-like operations over variables, free generalization comes naturally.

Furthermore, he specifies algebraic rules further and he writes:

Marcus (2001)

Algebraic rules are not finite tables of memorized facts or relationships between specific instances but open-ended relationships that can be freely generalized to all elements within some class.

Most decisively, for Marcus, a UQOTOM isn't just a special type of function; it is the very implementation of an algebraic rule. He writes:

Marcus (2001)

Premise 2.1.1. *When such a network represents identity or some other UQOTOM, it represents an abstract relationship between variables—which is to say that such a network implements an algebraic rule.*

Critique

The implication

$$f \text{ is a UQOTOM} \implies f \text{ implements an algebraic rule}$$

presents a foundational but significant oversimplification. By narrowing the scope of "algebraic rules" to only one-to-one mappings, this framework fails to account for the true complexity of algebraic operations and serves as a poor indicator of a model's reasoning capabilities.

For instance, a common operation like the quadratic function, $f(x) = x^2$, is undeniably an algebraic rule, yet it is not a UQOTOM because it is not one-to-one (e.g., $f(2) = f(-2) = 4$). The fundamental flaw in Marcus's approach is this reduction of the broad, complex class of algebraic rules to the narrow, simpler subset of UQOTOMs. A model's ability to handle this special case provides insufficient evidence of a general capacity for algebraic thought.

2.1.1 Humans can freely generalize UQOTOMs

Now that Marcus has established his framework of UQOTOMs, he argues that humans can freely generalize these UQOTOMs. He provides evidence for this claim with the following example:

Example 2.1.1. When you were presented this pattern

Input	Output
1010	1010
0100	0100
1110	1110
0000	0000
1111	?

and asked to predict the output of 1111, what would you answer?

Marcus argues that most people would answer 1111, because humans have a bias for generalizing UQOTOMs, which is why the identity mapping is generalized in this example. He follows from this empirical observation:

Marcus (2001) p. 50

Premise 2.1.2. *The bottom line is that humans can freely generalize one-to-one mappings [...]*

Thus, in Marcus's line of argumentation, humans can generalize algebraic rules based on the previous result and premise 2.1.1. He concludes:

Conclusion 2.1.1. *Models of human cognition must be able to generalize UQOTOMs, and thus algebraic rules.*

This framework of a necessary condition for models of human cognition is the basis for Marcus's following analysis.

2.1.2 Multilayer Perceptrons and Operations over Variables

Next, Marcus emphasizes the distinction between models that allocate *one node per variable*, and models which allocate *multiple nodes per variable*:

Definition 2.1.2. A model is said to allocate *one node per variable* if it has a single node for each variable in the input. A model is said to allocate *multiple nodes per variable* if it has multiple nodes for each variable in the input.

Marcus (2001)

Again, what is relevant here is not the sheer number of input units but rather the number of input units allocated to representing each input variable.

Critique

In the mathematical setting of neural networks that represent one-to-one mappings, it makes no difference whether, say, 5 input nodes all together represent a single variable, or five individual variables all encoded by a single node each. Furthermore, he also does not apply this framework, or only very vaguely, as we will see in the following.

Additionally, he points out that this distinction is not to be confused with *localist* and *distributed* representations:

Marcus (2001)

While all models that use distributed representations allocate more than one variable per node, it is not the case that all localist models allocate a single node per variable.

With this definition established, Marcus claims that:

Marcus (2001)

One-node-per-variable models, it turns out, can and indeed (the caveats in note 5 notwithstanding) must represent universally quantified one-to-one mappings.

Critique

Again, he assumes the network to only have one input node, which is not the same as having one node per variable. Furthermore, in note 5 he explains that he assumes the network to only have linear activation functions. Such a network must learn a linear mapping, and every mapping $f(x) = \alpha \cdot x$ is one-to-one for $\alpha \neq 0$.

Note that if we had multiple input nodes instead, a linear mapping is described by a matrix multiplication $f(\mathbf{x}) = A \cdot \mathbf{x}$. In this case, the mapping is one-to-one if and only if A is invertible, which is not guaranteed for all matrices.

Critique

This described case of a single input node with linear activation functions is very specific and doesn't serve any general argument. Note that when using non-linear activation functions, one can easily construct a network representing a non one-to-one mapping:

Example 2.1.2. MLPs with non-linear activations like $\tanh(x)$ can represent non-injective functions (other than $f(x) \equiv \mathbf{0}$) even when only using one input node. For instance, the MLP depicted in figure 2.1 implements the non-injective mapping depicted in figure 2.2.




Figure 2.1: Simple MLP using $\tanh(x)$ activation implementing a non-injective mapping.




Figure 2.2: Plot of the function $y = g(g(x) + g(-2x))$ with $g(x) := \tanh(x)$.

Regarding models that allocate more than one node per variable, Marcus states:

Marcus (2001)

Models that allocate more than one node per variable too, can represent universally quantified one-to-one mappings (see, for example, the left panel of figure 3.4), but they do not have to (see the right panel of figure 3.4).

Critique

This statement carries very little information. It only states that there exists MLPs that represent UQOTOMs, and that there exist MLPs that do not represent UQOTOMs, which is trivially true. The bigger problem, however, is that he also does not provide any further details on how to construct such MLPs, or what the implications of this distinction are.

Learning

Until now, Marcus has only discussed the ability of MLPs to represent UQOTOMs. Now, he turns to the question of whether MLPs can *learn* UQOTOMs. He states:

Marcus (2001)

The flexibility in what multiple-nodes-per-variable models can represent leads to a flexibility in what they can learn. Multiple-nodes-per-variable models can learn UQOTOMs, and they can learn arbitrary mappings. But what they learn depends on the nature of the learning algorithm. Back-propagation—the learning algorithm most commonly used—does not allocate special status to UQOTOMs. Instead, a many-nodes-per-variable multilayer perceptron that is trained by back-propagation can learn a UQOTOM—such as identity, multiplication, or concatenation—only if it sees that UQOTOM illustrated with respect to each possible input and output node.

He justifies that back-propagation does not allocate special status to UQOTOMs with an example:

Example 2.1.3. The autoencoder network depicted in 2.3 is trained to learn the identity mapping:




Figure 2.3: Autoencoder network trained to learn the identity mapping.

However, all training samples have a 0 as a last bit. Marcus's simulations show that the network learns to ignore the last bit by always predicting a 0, and only learns the identity mapping for the first three bits.

He formalizes this failure of back-propagation to learn UQOTOMs by defining the concept of *training independence*:

Marcus (2001)

The equations lead to two properties that I call input independence and output independence or, collectively, training independence (Marcus, 1998c).

Marcus (2001) p. 47

Definition 2.1.3 (Input Independence). "Input independence is about how the connections that emanate from input nodes are trained. First, when an input node is always off (that is, set to 0), the connections that emanate from it will never change. This is because the term in the equation that determines the size of the weight change for a given connection from input node x into the rest of the network is always multiplied by the activation of input node x ; if the activation of input node x is 0, the connection weight does not change.

Marcus (2001) p. 47

Definition 2.1.4 (Output Independence). Output independence is about the connections that feed into the output units. The equations that adjust the weights feeding an output unit j depend on the difference between the observed output for unit j and the target output for unit j but not on the observed values or target values of any other unit. Thus the way the network adjusts the weights feeding output node j must be independent of the way the network adjusts the weights feeding output node k (assuming that nodes j and k are distinct).

Critique

Marcus's concept of input independence is very specific and is not really applicable for real world examples. When an input node is always set to zero in training, we could simply use a different encoding scheme. Marcus argues that for an input node that is always set to the same constant v , the network "does not learn anything about the relation between that input node and the output, other than to always set the output node to value v " (Marcus 2001, p. 47).

This fact, however, is only based on empirical experiments, and is specific to the tested model architecture. What really happens is that the network won't use this node's value for inferring the output, but rather treats it as an adjustable bias for the layer it feeds into. This is not a flaw, but rather a property of the generic architectural design of MLPs.

Critique

Similarly, output independence is just another property of MLPs which isn't necessarily a flaw. It is not the case that the network learns to predict each output node based on the input nodes independently, since the hidden layers are trained based on *all* output nodes. Marcus acknowledges this fact, he writes:

"This means not that there is never any dependence between output nodes but that the only source of dependence between them is their common influence on the hidden nodes, which turns out not to be enough. At best, the mutual influence of output nodes on input-to-hidden- layer connections may under some circumstances lead to felicitous encodings of the hidden nodes. We might think of such hidden units as quasi-input nodes. The crucial point is that no matter how felicitous the choice of quasi-input units may be, the network must always learn the mapping between these quasi-input nodes and the output nodes. output nodes on input-to-hidden- layer connections is not sufficient to allow the network to generalize a UQOTOM between nodes." (Marcus 2001, p. 47 f.)

However, MLPs are not designed to generalize UQOTOMs, they are generic function approximators. Now, maybe this is exactly what Marcus criticizes, but how can we blame a generic model to not have an arbitrary bias, in this case that of UQOTOMs, when it is not designed for this purpose? When we want our model to have a specific bias, in this case to employ algebraic rules, we have to either learn the entire domain, or integrate domain specific knowledge into the architectural design (or else the model architecture wouldn't be generic).

2.2 Marcus's proposed Alternatives

He begins by stating five properties a successful model must fulfill:

Marcus (2001) p. 51

In general, what is required is a system that has five properties. First, the system must have a way to distinguish variables from instances, analogous to the way mathematics textbooks set variables in italic type (x) and constants in bold type (**AB**). Second, the system must have a way to represent abstract relationships between variables, analogous to an equation like $y = x + 2$. Third, the system must have a way to bind a particular instance to a given variable, just as the variable x may be assigned the value 7. Fourth, the system must have a way to apply operations to arbitrary instances of variables—for example, an addition operation must be able to take any two numbers as input, a copying operation must be able to copy any input, or a concatenation operation must be able to combine any two inputs. Finally, the system must have a way to extract relationships between variables on the basis of training examples.

Critique

These criteria are not really measurable. They seem to rather act as a guideline for designing a model architecture. For instance, how to measure whether a system has a way to distinguish variables from instances? Is model trained on both instances and variables, and what does it mean to treat the differently? How to measure whether a system has a way to extract relationships between variables on the basis of training examples? Does training data set cover the entire domain, or does the model have to generalize based on a small set of training examples, and how would we measure this generalization?

2.2.1 Encoding Schemes

Next in Marcus's line of argument, he presents different schemes to encode variables with input nodes. He wants to show that the inability of MLPs to learn UQOTOMs is fundamentally rooted in the architecture itself, and cannot easily be solved by different encoding schemes.

First, he introduces *conjunctive coding*.

Marcus (2001) p. 52

But the brain must rely on other techniques for variable binding as well. Conjunctive codes do not naturally allow for the representation of binding between a variable and a novel instance.

According to Marcus *tensor products* cannot solve all our problems as well:

Marcus (2001) p. 54

Nonetheless, despite these advantages, I suggest in chapter 4 that tensor products are not plausible as an account of how we represent recursively structured representations.

Lastly, he mentions *temporal synchrony*:

Marcus (2001) p. 57

My own view is that temporal synchrony might well play a role in some aspects of vision, such as grouping of parts of objects, but I have doubts about whether it plays as important a role in cognition and language.

Critique

It is a bit unclear why Marcus introduces these encoding schemes. He wants to show that the inability of MLPs to learn UQOTOMs is fundamentally rooted in the architecture itself, and cannot easily be solved by different encoding schemes. However, instead of providing a general argument for a generic class of encodings, he only discusses a few specific encodings.

Registers

Marcus suggests the use of *registers*. He writes:

Marcus (2001) p. 54

A limitation of the binding schemes discussed so far is that none provides a way of storing a binding. The bindings that are created are all entirely transitory, commonly taken to be constructed as the consequence of some current input to the system. One also needs a way to encode more permanently bindings between variables and instances.

Marcus (2001) p. 55

Registers are central to digital computers; my suggestion is that registers are central to human cognition as well.

Additionally, according to Marcus, registers could explain the phenomenon of *rapidly updatable memory*, which describes how we humans are able to learn things on a single trail, like remembering where we parked our car. The brain would use these registers to perform basic operations on them:

Marcus (2001) p. 58

My hunch is that the brain contains a similar stock of basic instructions, each defined to operate over all possible values of registers.

Marcus claims that this *modus operandi* would explain our ability to generalize algebraic rules:

Marcus (2001) p. 58

But in some cases we manage to extract a relationship between variables on the basis of training examples, without being given an explicit high-level description. Either way, when we learn a new function, we are presumably choosing among ways of combining some set of elementary instructions.

Critique

First, his claim that the human brain employs registers and basic operations on them is very bold, since he offers very little empirical data, relying instead on speculative arguments. Second, while Marcus gives some good arguments for his suggestion of using registers, it is still a bit unmotivated, especially with respect to the previous discussion. How do registers help a model to generalize UQOTOM? Marcus doesn't provide any explanation about such models and their behavior. The only purpose for his framework is to criticize the status quo, but he seems to have double standards for his own claims.

2.2.2 Case Studies

To substantiate his theoretical framework, Marcus concludes the chapter with two detailed case studies: the learning of artificial grammars by infants and the inflection of the English past tense. In both examples, he argues that models incorporating rule-like, symbolic operations—similar to the register-based architecture he advocates—can successfully account for human cognitive phenomena where he claims standard connectionist models fail.

A full analysis of these specific empirical arguments is beyond the scope of this paper, which has focused on the chapter's core theoretical claims and their weaknesses. However, these case studies represent Marcus's primary evidence for the necessity of the architectural properties he proposes.

2.3 Summary

In this chapter of Marcus's book *The Algebraic Mind*, the author established a new framework for assessing the capability of models to learn algebraic rules.

Marcus supports his claims with examples and empirical data. Nonetheless, the reader should remain critical of his framework and the arguments given. Central points of criticism are Marcus's reduction of algebraic rules to what he calls UQOTOMs, the consequences of his critique of backpropagation, and lastly his claims about registers and the human brain, which can be seen as unmotivated and bold.

While Marcus's approach is arguably too simplified for the complex nature of symbol manipulation and human cognition, his critique should not be overlooked. Even though his framework may lack some formal rigor, his main points—namely, the apparent misalignment between human cognitive biases and those of popular architectures like MLPs, as well as his appeal to look for alternatives—remain relevant today.

Meta question: Learning vs. inference question -; he often refers to human capabilities: "Cases in which humans can freely generalize UQOTOM on the basis of restricted data are problematic for multiple-nodes-per-variable multi-layer perceptrons trained by back-propagation." (Marcus 2001, p. 50)

Summary:

Hit critique is valid, but it is obvious and explanations are too simple to infer any meaningful conclusions. - too simple formalization of complex concept

The claims he made at the end are very bold and not really well argued. One can either believe them or not.

3 Contemporary Research

Even though Marcus published his book in 2001, he still maintains his main points of criticism regarding popular connectionist models. In his blog posts he discusses concurrent research, and he points out fundamental weaknesses of popular connectionist models, including LLMs.

Marcus (2025a)

LLM are giant, opaque black boxes with no explicit models of the world at all. Part of what it means to say that an LLM is a black box is to say that you can't point to an articulated model of any particular set of facts inside.

Marcus (2025a)

But the LLM doesn't maintain structured symbolic systems like databases; it has no direct database of cities and populations. Aside from what it can retrieve with external tools like web searches (which provide a partial but only partial workaround) it just has a bunch of bits of text that it stochastically reconstructs. And that is why they hallucinate, frequently. They wouldn't need to do so if they actually had reliable access to relevant databases.

3.1 The Illusion of Thinking

In order to improve the symbolic capabilities of LLMs, researchers suggested various methods, including *Large Reasoning Models* (LRMs). Shojae^{*†} et al. (2025) acknowledge this in their paper *The Illusion of Thinking*:

While these LLMs demonstrate promising language understanding with strong compression capabilities, their intelligence and reasoning abilities remain a critical topic of scientific debate [7, 8]. Earlier iterations of LLMs [9, 10, 11] exhibited poor performance on reasoning benchmarks [12, 13, 14, 6]. To address these shortcomings, several approaches have been explored with the common theme among them being “scaling” both the training data and test-time computation. For instance, generating a Chain of Thought (CoT) [15, 16, 17, 18] and incorporating self-verification [19, 20, 21] prior to the final answer have been shown to improve model performance.

The authors of this paper analyzed how well standard LLMs could solve classical search problems like *The Tower of Hanoi*, and to what extent LRMs improve symbolic capabilities. They tested several models on varying problem complexities, like different number of starting disks in the case of Tower of Hanoi.

What they found is that both LLMs and LRMs have high accuracy on simple problems. As complexity increases, reasoning models outperform their non-reasoning counterparts, until both reach a point where the models are unable to solve the problem at all. These findings are summarized in figure 3.1.




Figure 3.1: Performance of LLMs and LRMs on Tower of Hanoi problems with varying complexity. Source: Shojaee*† et al. (2025).

Other conclusions of the paper are:

Shojaee*† et al. (2025)

Our detailed analysis of reasoning traces further exposed complexity-dependent reasoning patterns, from inefficient “overthinking” on simpler problems to complete failure on complex ones. These insights challenge prevailing assumptions about LRM capabilities and suggest that current approaches may be encountering fundamental barriers to generalizable reasoning. Finally, we presented some surprising results on LRMs that lead to several open questions for future work. Most notably, we observed their limitations in performing exact computation; for example, when we provided the solution algorithm for the Tower of Hanoi to the models, their performance on this puzzle did not improve. Moreover, investigating the first failure move of the models revealed surprising behaviors. For instance, they could perform up to 100 correct moves in the Tower of Hanoi but fail to provide more than 5 correct moves in the River Crossing puzzle.

Marcus argues that this paper provides evidence for his claim of poor symbolic capabilities of connectionist models due to insufficient *structured symbolic systems*:

Marcus (2025b)

The new Apple paper adds to the force of Rao’s critique (and my own) by showing that even the latest of these new-fangled “reasoning models” still —even having scaled beyond o1— fail to reason beyond the distribution reliably, on a whole bunch of classic problems, like the Tower of Hanoi. For anyone hoping that “reasoning” or “inference time compute” would get LLMs back on track, and take away the pain of multiple failures at getting pure scaling to yield something worthy of the name GPT-5, this is bad news.

Critique

While the findings of Shojaee*† et al. (2025) are compelling, their methodology warrants critical examination. As the authors themselves acknowledge, the reliance on highly algorithmic tasks like the Tower of Hanoi may not be representative of all forms of reasoning. Furthermore, given the known sensitivity of LLMs to prompt phrasing, it remains an open question whether the observed performance limits are fundamental to the models themselves or an artifact of the specific prompting strategies employed.

Additionally, Marcus’s conclusion that these shortcomings prove the models are incapable of reasoning is too strong. The paper itself shows that the models have better reasoning capabilities on problems that are more popular, suggesting performance is heavily influenced by representation in the training data. While an untrained model clearly lacks reasoning abilities, these skills demonstrably improve with training. Therefore, one cannot rule out the possibility that with sufficient scale and the right data, these models could eventually meet the demands of complex symbolic tasks.

3.2 A Case for Neuro-Symbolic AI: AlphaGeometry

After having analyzed Marcus’s critique of purely connectionist models, it seems less surprising that he is an advocate of *neuro-symbolic AI*, which integrates neural and symbolic AI architectures to address the weaknesses of each. Marcus writes in his article “AlphaProof, AlphaGeometry, ChatGPT, and why the future of AI is neurosymbolic”:

Marcus (2024)

The idea is to try to take the best of two worlds, combining (akin to Kahneman’s System I and System II), neural networks, which are good at kind of quick intuition from familiar examples (a la Kahneman’s System I) with explicit symbolic systems that use formal logic and other reasoning tools (a la Kahneman’s System II).

One of these hybrid models is Google’s *alphageometry*, which the authors described in “Solving olympiad geometry without human demonstrations”. Amongst all the hype around connectionist models, Marcus highly values this paper:

Marcus (2024)

Which brings me to one of the few AI reports I have liked lately, their blog this week on progress on the International Math Olympiad, in which (with some caveats) they achieved Silver Medal performance, far ahead of what most humans could manage.

In order to understand the claimed interplay of connectionism and symbolism, we have to take a closer look at the paper. These are the results in a nutshell according to the authors:

Trinh et al. (2024)

AlphaGeometry is a neuro-symbolic system that uses a neural language model, trained from scratch on our large-scale synthetic data, to guide a symbolic deduction engine through infinite branching points in challenging problems. On a test set of 30 latest olympiad-level problems, AlphaGeometry solves 25, outperforming the previous best method that only solves ten problems and approaching the performance of an average International Mathematical Olympiad (IMO) gold medallist.




Figure 3.2: Overview of the neuro-symbolic AlphaGeometry model.

Source: Trinh et al. (2024)

One can think of the model this way: The deduction engine explores a search graph, where it is given the starting point, that is the premises of the problem, and the end-point, that is the theorem to prove. Based on *deductive database*, the symbolic engine can explore paths in this network. However, since some problems require additional assumptions to be made, like defining new points, an exhaustive search may not always yield the answer right away.

This is where language model comes in. It tries to predict the necessary steps needed to take in the proof, where its intuition serves to create new premises. This information is then fed back into the deduction engine. This cycle continues until the engine infers the desired statement of interest, or until a timeout occurs.

In the same blog post where Marcus discusses these advancements, he concludes:

Marcus (2024)

There can be no path to AGI without neurosymbolic AI.

This conclusion is the result of his persistent critique on purely connectionist models, as well as results like AlphaGeometry which mark great advancements in AI and are based on neuro-symbolic architectures.

4 Summary

For over two decades, Gary Marcus has established himself as a prominent and persistent critic of purely connectionist approaches to artificial intelligence. His 2001 book, *The Algebraic Mind*, laid the foundation for his main line of argumentation: that connectionist models, on their own, lack the necessary architecture for true generalization and the kind of systematic, algebraic reasoning that is central to human cognition.

As this analysis has shown, Marcus’s specific frameworks—such as his reduction of algebraic rules to UQOTOMs—and the formal structure of his arguments can be challenged for their simplicity and speculative nature. His proposed solutions, like a register-based architecture, often lack the concrete, testable detail he demands of the models he critiques.

Despite these methodological weaknesses, his core message should not be overlooked. Marcus’s fundamental critique—that a field focused solely on one paradigm risks ignoring essential components of intelligence—remains highly relevant. Even today, he actively engages with contemporary research, interpreting findings from papers on both the limitations of LLMs and the successes of neuro-symbolic systems like AlphaGeometry as further evidence for his long-standing claims.

Now that we have presented Marcus’s views, the empirical evidence he draws upon, and the critiques of his own framework, it is up to the reader to decide on which points they agree with Marcus or where the weaknesses in his arguments lie. We hope to have illuminated this important debate from multiple viewpoints, so that the reader may form their own critical opinion.

Bibliography

- Marcus, G. F. (2001). *The algebraic mind: Integrating connectionism and cognitive science*. The MIT Press.
- Marcus, G. F. (2024, July). *Alphaproof, alphageometry, chatgpt, and why the future of ai is neurosymbolic*. Retrieved August 7, 2025, from <https://garymarcus.substack.com/p/alphaproof-alphageometry-chatgpt>
- Marcus, G. F. (2025a, June). *Generative ai's crippling and widespread failure to induce robust models of the world*. Retrieved August 7, 2025, from <https://garymarcus.substack.com/p/generative-ais-crippling-and-widespread>
- Marcus, G. F. (2025b, June). *A knockout blow for llms?* Retrieved August 7, 2025, from <https://garymarcus.substack.com/p/a-knockout-blow-for-llms>
- Shojaee*†, P., Mirzadeh*, I., Alizadeh, K., Horton, M., Bengio, S., & Farajtabar, M. (2025). The illusion of thinking: Understanding the strengths and limitations of reasoning models via the lens of problem complexity. <https://ml-site.cdn-apple.com/papers/the-illusion-of-thinking.pdf>
- Trinh, T. H., Wu, Y., Le, Q. V., He, H., & Luong, T. (2024). Solving olympiad geometry without human demonstrations. *Nature*, 625(7995), 476–482. <https://doi.org/10.1038/s41586-023-06747-5>